

Muhammad Bilal Iqbal

July 5, 2020

1 Data Cleaning with Python

It's commonly said that data scientists spend 80% of their time cleaning and manipulating data and only 20% of their time analyzing it. The time spent cleaning is vital since analyzing dirty data can lead us to draw inaccurate conclusions. Data cleaning is an essential task in data science. Without properly cleaned data, the results of any data analysis or machine learning model could be inaccurate. Here I will work to identify, diagnose, and treat a variety of data cleaning problems in Python, ranging from simple to advanced. I will deal with improper data types, check that the data is in the correct range, handle missing data, perform record linkage, and more!

Common Data Problems

Here I'll learn how to overcome some of the most common dirty data problems. I'll convert data types, apply range constraints to remove future data points, and remove duplicated data points to avoid double-counting.

Common data types

Manipulating and analyzing data with incorrect data types could lead to compromised analysis as you go along the data science workflow.

When working with new data, you should always check the data types of your columns using the `.dtypes` attribute or the `.info()` method which I'll work on next. Often times, I'll run into columns that should be converted to different data types before starting any analysis.

Correctly identifying what type your data is is one of the easiest ways to avoid hampering your analysis due to data type constraints in the long run.

Numeric data or ... ?

I'll be working with bicycle ride sharing data in San Francisco called `ride_sharing`. It contains information on the start and end stations, the trip duration, and some user information for a bike sharing service.

The `user_type` column contains information on whether a user is taking a free ride and takes on the following values:

- 1 - for free riders.
- 2 - for pay per ride.
- 3 - for monthly subscribers.

I will print the information of ride_sharing using .info() and see a firsthand example of how an incorrect data type can flaw the analysis of the dataset. The pandas package is imported as pd.

```
[1]: import pandas as pd
```

```
[4]: ride_sharing = pd.read_csv('ride_sharing_new.csv', index_col=0)
```

```
[7]: # Printing the information of ride_sharing
print(ride_sharing.info())

# Printing summary statistics of user_type column
print(ride_sharing['user_type'].describe())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25760 entries, 0 to 25759
Data columns (total 9 columns):
duration                25760 non-null object
station_A_id            25760 non-null int64
station_A_name          25760 non-null object
station_B_id            25760 non-null int64
station_B_name          25760 non-null object
bike_id                 25760 non-null int64
user_type               25760 non-null int64
user_birth_year         25760 non-null int64
user_gender             25760 non-null object
dtypes: int64(5), object(4)
memory usage: 2.0+ MB
None
count    25760.000000
mean      2.008385
std       0.704541
min       1.000000
25%       2.000000
50%       2.000000
75%       3.000000
max       3.000000
Name: user_type, dtype: float64
```

```
[6]: ride_sharing.head()
```

```
[6]:      duration  station_A_id \
0  12 minutes         81
1  24 minutes          3
2   8 minutes        67
3   4 minutes        16
4  11 minutes        22
```

```
station_A_name  station_B_id \
```

0		Berry St at 4th St	323
1		Powell St BART Station (Market St at 4th St)	118
2		San Francisco Caltrain Station 2 (Townsend St...	23
3		Steuart St at Market St	28
4		Howard St at Beale St	350

	station_B_name	bike_id	user_type	user_birth_year	\
0	Broadway at Kearny	5480	2	1959	
1	Eureka Valley Recreation Center	5193	2	1965	
2	The Embarcadero at Steuart St	3652	3	1993	
3	The Embarcadero at Bryant St	1883	1	1979	
4	8th St at Brannan St	4626	2	1994	

	user_gender
0	Male
1	Male
2	Male
3	Male
4	Male

Question

By looking at the summary statistics - they don't really seem to offer much description on how users are distributed along their purchase type, why do you think that is?

Answer:

The `user_type` column has a finite set of possible values that represent groupings of data, it should be converted to category.

Convert `user_type` into categorical by assigning it the 'category' data type and store it in the `user_type_cat` column.

Make sure you converted `user_type_cat` correctly by using an assert statement.

```
[8]: # Printing the information of ride_sharing
print(ride_sharing.info())

# Printing summary statistics of user_type column
print(ride_sharing['user_type'].describe())

# Converting user_type from integer to category
ride_sharing['user_type_cat'] = ride_sharing['user_type'].astype('category')

# Writing an assert statement confirming the change
assert ride_sharing['user_type_cat'].dtype == 'category'

# Printing new summary statistics
print(ride_sharing['user_type_cat'].describe())
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 25760 entries, 0 to 25759
Data columns (total 9 columns):
duration                25760 non-null object
station_A_id            25760 non-null int64
station_A_name          25760 non-null object
station_B_id            25760 non-null int64
station_B_name          25760 non-null object
bike_id                 25760 non-null int64
user_type               25760 non-null int64
user_birth_year         25760 non-null int64
user_gender             25760 non-null object
dtypes: int64(5), object(4)
memory usage: 2.0+ MB
None
count      25760.000000
mean         2.008385
std          0.704541
min          1.000000
25%          2.000000
50%          2.000000
75%          3.000000
max          3.000000
Name: user_type, dtype: float64
count      25760
unique       3
top          2
freq       12972
Name: user_type_cat, dtype: int64

```

```
[9]: ride_sharing.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 25760 entries, 0 to 25759
Data columns (total 10 columns):
duration                25760 non-null object
station_A_id            25760 non-null int64
station_A_name          25760 non-null object
station_B_id            25760 non-null int64
station_B_name          25760 non-null object
bike_id                 25760 non-null int64
user_type               25760 non-null int64
user_birth_year         25760 non-null int64
user_gender             25760 non-null object
user_type_cat           25760 non-null category
dtypes: category(1), int64(5), object(4)
memory usage: 2.0+ MB

```

```
[ ]: Taking a look at the new summary statistics, it seems that most users are pay_u
    ↳per ride users!
```

```
[10]: # Stripping duration of minutes
ride_sharing['duration_trim'] = ride_sharing['duration'].str.strip('minutes')

# Converting duration to integer
ride_sharing['duration_time'] = ride_sharing['duration_trim'].astype('int')

# Writing an assert statement making sure of conversion
assert ride_sharing['duration_time'].dtype == 'int'

# Printing formed columns and calculate average ride duration
print(ride_sharing[['duration', 'duration_trim', 'duration_time']])
print(ride_sharing['duration_time'].mean())
```

	duration	duration_trim	duration_time
0	12 minutes	12	12
1	24 minutes	24	24
2	8 minutes	8	8
3	4 minutes	4	4
4	11 minutes	11	11
...
25755	11 minutes	11	11
25756	10 minutes	10	10
25757	14 minutes	14	14
25758	14 minutes	14	14
25759	29 minutes	29	29

```
[25760 rows x 3 columns]
11.389052795031056
```

11 minutes is really not bad for an average ride duration in a city like San-Francisco. Next I am going to jump right ahead into sanity checking the range of values in the data.

```
[11]: ride_sharing.head()
```

```
[11]:      duration  station_A_id \
0  12 minutes           81
1  24 minutes            3
2   8 minutes          67
3   4 minutes          16
4  11 minutes          22
```

	station_A_name	station_B_id \
0	Berry St at 4th St	323
1	Powell St BART Station (Market St at 4th St)	118
2	San Francisco Caltrain Station 2 (Townsend St...	23

3	Steuart St at Market St	28
4	Howard St at Beale St	350

	station_B_name	bike_id	user_type	user_birth_year	\
0	Broadway at Kearny	5480	2	1959	
1	Eureka Valley Recreation Center	5193	2	1965	
2	The Embarcadero at Steuart St	3652	3	1993	
3	The Embarcadero at Bryant St	1883	1	1979	
4	8th St at Brannan St	4626	2	1994	

	user_gender	user_type_cat	duration_trim	duration_time
0	Male	2	12	12
1	Male	2	24	24
2	Male	3	8	8
3	Male	1	4	4
4	Male	2	11	11

Tire size constraints

Here I am going to build on top of the work you've been doing with the `ride_sharing` DataFrame. I'll be working with the `tire_sizes` column which contains data on each bike's tire size.

Bicycle tire sizes could be either 26 , 27 or 29 and are here correctly stored as a categorical value. In an effort to cut maintenance costs, the ride sharing provider decided to set the maximum tire size to be 27 .

I will make sure the `tire_sizes` column has the correct range by first converting it to an integer, then setting and testing the new upper limit of 27 for tire sizes.

```
[ ]: # Convert tire_sizes to integer
#ride_sharing['tire_sizes'] = ride_sharing['tire_sizes'].astype('int')

# Set all values above 27 to 27
#ride_sharing.loc[ride_sharing['tire_sizes'] > 27, 'tire_sizes'] = 27

# Reconvert tire_sizes back to categorical
#ride_sharing['tire_sizes'] = ride_sharing['tire_sizes'].astype('category')

# Print tire size description
#print(ride_sharing['tire_sizes'].describe())
```

Now I will work on another airlines dataset and deal with text and categorical data problems

```
[49]: import datetime as dt
today_date = dt.date.today()
```

```
[50]: today_date
```

```
[50]: datetime.date(2020, 7, 5)
```

```
[17]: airlines = pd.read_csv('airlines_final.csv')
airlines.head()
```

```
[17]:
```

	Unnamed: 0	id	day	airline	destination	dest_region \
0	0	1351	Tuesday	UNITED INTL	KANSAI	Asia
1	1	373	Friday	ALASKA	SAN JOSE DEL CABO	Canada/Mexico
2	2	2820	Thursday	DELTA	LOS ANGELES	West US
3	3	1157	Tuesday	SOUTHWEST	LOS ANGELES	West US
4	4	2992	Wednesday	AMERICAN	MIAMI	East US

	dest_size	boarding_area	dept_time	wait_min	cleanliness \
0	Hub	Gates 91-102	2018-12-31	115.0	Clean
1	Small	Gates 50-59	2018-12-31	135.0	Clean
2	Hub	Gates 40-48	2018-12-31	70.0	Average
3	Hub	Gates 20-39	2018-12-31	190.0	Clean
4	Hub	Gates 50-59	2018-12-31	559.0	Somewhat clean

	safety	satisfaction
0	Neutral	Very satisfied
1	Very safe	Very satisfied
2	Somewhat safe	Neutral
3	Very safe	Somewhat satisfied
4	Very safe	Somewhat satisfied

Text and categorical data problems

Categorical and text data can often be some of the messiest parts of a dataset due to their unstructured nature. Here I will work to fix whitespace and capitalization inconsistencies in category labels, collapse multiple categories into one, and reformat strings for consistency.

```
[51]: categories = pd.read_csv('survey.csv')
categories
```

```
[51]:
```

	cleanliness	safety	satisfaction
0	Clean	neutral	Very satisfied
1	Average	Very safe	Neutral
2	Somewhat clean	Somewhat safe	Somewhat satisfied
3	Somewhat dirty	Very unsafe	Somewhat unsatisfied
4	Dirty	Somewhat unsafe	Very unsatisfied

```
[18]: # Print categories DataFrame
print(categories)

# Print unique values of survey columns in airlines
print('Cleanliness: ', airlines['cleanliness'].unique(), "\n")
print('Safety: ', airlines['safety'].unique(), "\n")
print('Satisfaction: ', airlines['satisfaction'].unique(), "\n")
```

	cleanliness	safety	satisfaction
0	Clean	neutral	Very satisfied
1	Average	Very safe	Neutral
2	Somewhat clean	Somewhat safe	Somewhat satisfied
3	Somewhat dirty	Very unsafe	Somewhat unsatisfied
4	Dirty	Somewhat unsafe	Very unsatisfied

Cleanliness: ['Clean' 'Average' 'Somewhat clean' 'Somewhat dirty' 'Dirty']

Safety: ['Neutral' 'Very safe' 'Somewhat safe' 'Very unsafe' 'Somewhat unsafe']

Satisfaction: ['Very satisfied' 'Neutral' 'Somewhat satisfied' 'Somewhat unsatisfied' 'Very unsatisfied']

```
[19]: # Finding the cleanliness category in airlines not in categories
cat_clean = set(airlines['cleanliness']).difference(categories['cleanliness'])

# Finding rows with that category
cat_clean_rows = airlines['cleanliness'].isin(cat_clean)

# Printing rows with inconsistent category
print(airlines[cat_clean_rows])
```

	Unnamed: 0	id	day	airline	destination	\
0	0	1351	Tuesday	UNITED INTL	KANSAI	
1	1	373	Friday	ALASKA	SAN JOSE DEL CABO	
3	3	1157	Tuesday	SOUTHWEST	LOS ANGELES	
6	6	2578	Saturday	JETBLUE	LONG BEACH	
10	11	1129	Tuesday	SOUTHWEST	SAN DIEGO	
...	
2470	2802	394	Friday	ALASKA	LOS ANGELES	
2473	2805	2222	Thursday	SOUTHWEST	PHOENIX	
2474	2806	2684	Friday	UNITED	ORLANDO	
2475	2807	2549	Tuesday	JETBLUE	LONG BEACH	
2476	2808	2162	Saturday	CHINA EASTERN	QINGDAO	

	dest_region	dest_size	boarding_area	dept_time	wait_min	cleanliness	\
0	Asia	Hub	Gates 91-102	2018-12-31	115.0	Clean	
1	Canada/Mexico	Small	Gates 50-59	2018-12-31	135.0	Clean	
3	West US	Hub	Gates 20-39	2018-12-31	190.0	Clean	
6	West US	Small	Gates 1-12	2018-12-31	63.0	Clean	
10	West US	Medium	Gates 20-39	2018-12-31	540.0	Clean	
...		
2470	West US	Hub	Gates 50-59	2018-12-31	115.0	Clean	
2473	West US	Hub	Gates 20-39	2018-12-31	165.0	Clean	
2474	East US	Hub	Gates 70-90	2018-12-31	92.0	Clean	
2475	West US	Small	Gates 1-12	2018-12-31	95.0	Clean	

2476	Asia	Large	Gates 1-12	2018-12-31	220.0	Clean
------	------	-------	------------	------------	-------	-------

	safety	satisfaction
0	Neutral	Very satisfied
1	Very safe	Very satisfied
3	Very safe	Somewhat satisfied
6	Very safe	Somewhat satisfied
10	Very safe	Somewhat satisfied
...
2470	Very safe	Very satisfied
2473	Very safe	Very satisfied
2474	Very safe	Very satisfied
2475	Somewhat safe	Very satisfied
2476	Very safe	Somewhat satisfied

[911 rows x 13 columns]

```
[20]: # Finding the cleanliness category in airlines not in categories
cat_clean = set(airlines['cleanliness']).difference(categories['cleanliness'])

# Finding rows with that category
cat_clean_rows = airlines['cleanliness'].isin(cat_clean)

# Printing rows with inconsistent category
print(airlines[cat_clean_rows])

# Printing rows with consistent categories only
print(airlines[~cat_clean_rows])
```

	Unnamed: 0	id	day	airline	destination \
0	0	1351	Tuesday	UNITED INTL	KANSAI
1	1	373	Friday	ALASKA	SAN JOSE DEL CABO
3	3	1157	Tuesday	SOUTHWEST	LOS ANGELES
6	6	2578	Saturday	JETBLUE	LONG BEACH
10	11	1129	Tuesday	SOUTHWEST	SAN DIEGO
...
2470	2802	394	Friday	ALASKA	LOS ANGELES
2473	2805	2222	Thursday	SOUTHWEST	PHOENIX
2474	2806	2684	Friday	UNITED	ORLANDO
2475	2807	2549	Tuesday	JETBLUE	LONG BEACH
2476	2808	2162	Saturday	CHINA EASTERN	QINGDAO

	dest_region	dest_size	boarding_area	dept_time	wait_min	cleanliness \
0	Asia	Hub	Gates 91-102	2018-12-31	115.0	Clean
1	Canada/Mexico	Small	Gates 50-59	2018-12-31	135.0	Clean
3	West US	Hub	Gates 20-39	2018-12-31	190.0	Clean
6	West US	Small	Gates 1-12	2018-12-31	63.0	Clean
10	West US	Medium	Gates 20-39	2018-12-31	540.0	Clean

...
2470	West US	Hub	Gates 50-59	2018-12-31	115.0	Clean
2473	West US	Hub	Gates 20-39	2018-12-31	165.0	Clean
2474	East US	Hub	Gates 70-90	2018-12-31	92.0	Clean
2475	West US	Small	Gates 1-12	2018-12-31	95.0	Clean
2476	Asia	Large	Gates 1-12	2018-12-31	220.0	Clean

	safety	satisfaction
0	Neutral	Very satisfied
1	Very safe	Very satisfied
3	Very safe	Somewhat satisfied
6	Very safe	Somewhat satisfied
10	Very safe	Somewhat satisfied
...
2470	Very safe	Very satisfied
2473	Very safe	Very satisfied
2474	Very safe	Very satisfied
2475	Somewhat safe	Very satisfied
2476	Very safe	Somewhat satisfied

[911 rows x 13 columns]

	Unnamed: 0	id	day	airline	destination \
2	2	2820	Thursday	DELTA	LOS ANGELES
4	4	2992	Wednesday	AMERICAN	MIAMI
5	5	634	Thursday	ALASKA	NEWARK
7	8	2592	Saturday	AEROMEXICO	MEXICO CITY
8	9	919	Friday	AIR CANADA	TORONTO
...
2466	2798	3099	Sunday	ALASKA	NEWARK
2468	2800	1942	Tuesday	UNITED	BOSTON
2469	2801	2130	Thursday	CATHAY PACIFIC	HONG KONG
2471	2803	2888	Wednesday	UNITED	AUSTIN
2472	2804	1475	Tuesday	ALASKA	NEW YORK-JFK

	dest_region	dest_size	boarding_area	dept_time	wait_min \
2	West US	Hub	Gates 40-48	2018-12-31	70.0
4	East US	Hub	Gates 50-59	2018-12-31	559.0
5	East US	Hub	Gates 50-59	2018-12-31	140.0
7	Canada/Mexico	Hub	Gates 1-12	2018-12-31	215.0
8	Canada/Mexico	Hub	Gates 91-102	2018-12-31	70.0
...
2466	East US	Hub	Gates 50-59	2018-12-31	105.0
2468	EAST US	Large	Gates 70-90	2018-12-31	145.0
2469	Asia	Hub	Gates 1-12	2018-12-31	380.0
2471	Midwest US	Medium	Gates 70-90	2018-12-31	60.0
2472	East US	Hub	Gates 50-59	2018-12-31	280.0

cleanliness	safety	satisfaction
-------------	--------	--------------

2	Average	Somewhat safe	Neutral
4	Somewhat clean	Very safe	Somewhat satisfied
5	Somewhat clean	Very safe	Very satisfied
7	Somewhat clean	Very safe	Neutral
8	Somewhat clean	Somewhat safe	Somewhat satisfied
...
2466	Somewhat clean	Very safe	Somewhat satisfied
2468	Somewhat clean	Somewhat safe	Somewhat satisfied
2469	Somewhat clean	Somewhat safe	Somewhat satisfied
2471	Somewhat clean	Somewhat unsafe	Somewhat satisfied
2472	Somewhat clean	Neutral	Somewhat satisfied

[1566 rows x 13 columns]

Inconsistent categories

I will be revisiting the airlines DataFrame from above.

As a reminder, the DataFrame contains flight metadata such as the airline, the destination, waiting times as well as answers to key questions regarding cleanliness, safety, and satisfaction on the San Francisco Airport.

I will examine two categorical columns from this DataFrame, `dest_region` and `dest_size` respectively, assess how to address them and make sure that they are cleaned and ready for analysis. The pandas package has been imported as `pd`, and the airlines DataFrame is in the environment.

```
[21]: # Printing the unique values in dest_region and dest_size respectively.
# Printing unique values of both columns
print(airlines['dest_region'].unique())
print(airlines['dest_size'].unique())
```

```
['Asia' 'Canada/Mexico' 'West US' 'East US' 'Midwest US' 'EAST US'
 'Middle East' 'Europe' 'eur' 'Central/South America'
 'Australia/New Zealand' 'middle east']
['Hub' 'Small' ' ' 'Hub' 'Medium' 'Large' 'Hub' ' ' 'Small'
 'Medium' ' ' 'Medium' 'Small' ' ' 'Large' 'Large' '']
```

```
[23]: # Changing the capitalization of all values of dest_region to lowercase.
# Replacing the 'eur' with 'europe' in dest_region using the .replace() method.
# Printing unique values of both columns
print(airlines['dest_region'].unique())
print(airlines['dest_size'].unique())

# Lowering dest_region column and then replace "eur" with "europe"
airlines['dest_region'] = airlines['dest_region'].str.lower()
airlines['dest_region'] = airlines['dest_region'].replace({'eur': 'europe'})
```

```
['asia' 'canada/mexico' 'west us' 'east us' 'midwest us' 'middle east'
 'europe' 'central/south america' 'australia/new zealand']
['Hub' 'Small' 'Medium' 'Large']
```

Stripping white spaces from the `dest_size` column using the `.strip()` method.

Verifying that the changes have been into effect by printing the unique values of the columns using `.unique()` .

```
[ ]: # Printing unique values of both columns
print(airlines['dest_region'].unique())
print(airlines['dest_size'].unique())

# Lowering dest_region column and then replace "eur" with "europe"
airlines['dest_region'] = airlines['dest_region'].str.lower()
airlines['dest_region'] = airlines['dest_region'].replace({'eur': 'europe'})

# Removing white spaces from `dest_size`
airlines['dest_size'] = airlines['dest_size'].str.strip()

# Verifying changes have been effected
print(airlines['dest_region'].unique())
print(airlines['dest_size'].unique())
```

Remapping categories

To better understand survey respondents from airlines, I want to find out if there is a relationship between certain responses and the day of the week and wait time at the gate.

The airlines DataFrame contains the `day` and `wait_min` columns, which are categorical and numerical respectively. The `day` column contains the exact day a flight took place, and `wait_min` contains the amount of minutes it took travelers to wait at the gate. To make the analysis easier, I will create two new categorical variables:

`wait_type`: 'short' for 0-60 min, 'medium' for 60-180 and long for 180+

`day_week`: 'weekday' if day is in the weekday, 'weekend' if day is in the weekend.

```
[25]: # Creating ranges for categories
import numpy as np
label_ranges = [0, 60, 180, np.inf]
label_names = ['short', 'medium', 'long']

# Creating wait_type column
airlines['wait_type'] = pd.cut(airlines['wait_min'], bins = label_ranges,
                              labels = label_names)

# Creating mappings and replace
mappings = {'Monday': 'weekday', 'Tuesday': 'weekday', 'Wednesday': 'weekday',
            'Thursday': 'weekday', 'Friday': 'weekday',
            'Saturday': 'weekend', 'Sunday': 'weekend'}

airlines['day_week'] = airlines['day'].replace(mappings)
```

```
[26]: airlines.head()
```

```
[26]:   Unnamed: 0   id   day   airline   destination   dest_region \
0          0  1351  Tuesday  UNITED INTL           KANSAI           asia
1          1   373   Friday    ALASKA  SAN JOSE DEL CABO  canada/mexico
2          2  2820  Thursday    DELTA    LOS ANGELES           west us
3          3  1157   Tuesday  SOUTHWEST    LOS ANGELES           west us
4          4  2992  Wednesday   AMERICAN           MIAMI           east us

   dest_size boarding_area  dept_time  wait_min  cleanliness \
0        Hub    Gates 91-102  2018-12-31    115.0         Clean
1      Small    Gates 50-59  2018-12-31    135.0         Clean
2        Hub    Gates 40-48  2018-12-31     70.0       Average
3        Hub    Gates 20-39  2018-12-31    190.0         Clean
4        Hub    Gates 50-59  2018-12-31    559.0  Somewhat clean

   safety      satisfaction  wait_type  day_week
0   Neutral    Very satisfied    medium  weekday
1   Very safe    Very satisfied    medium  weekday
2  Somewhat safe      Neutral    medium  weekday
3   Very safe  Somewhat satsified    long  weekday
4   Very safe  Somewhat satsified    long  weekday
```

Advanced data problems

I will dive into more advanced data cleaning problems, such as ensuring that weights are all written in kilograms instead of pounds. This analysis will help us to verify that values have been added correctly and that missing values don't negatively impact the analyses.

Ambiguous dates

I have a DataFrame containing a `subscription_date` column that was collected from various sources with different Date formats such as YYYY-mm-dd and YYYY-dd-mm. What is the best way to unify the formats for ambiguous values such as 2019-04-07?

- Set them to NA and drop them.
- Infer the format of the data in question by checking the format of subsequent and previous values.
- Infer the format from the original data source.
- All of the above are possible, as long as we investigate where our data comes from, and understand the dynamics affecting it before cleaning it.

Like most cleaning data tasks, ambiguous dates require a thorough understanding of where the data comes from. Diagnosing problems is the first step in finding the best solution!

Uniform currencies

I will now be working with a retail banking dataset stored in the banking DataFrame. The dataset contains data on the amount of money stored in accounts, their currency, amount invested, account opening date and last transaction date that were consolidated from American and European

branches.

Here I will be working with understanding the average account size and how investments vary by the size of account, however in order to produce this analysis accurately, we first need to unify the currency amount into dollars.

```
[27]: banking = pd.read_csv('banking_dirty.csv')
```

```
[28]: banking.head()
```

```
[28]:
```

	Unnamed: 0	cust_id	birth_date	Age	acct_amount	inv_amount	fund_A	\
0	0	870A9281	1962-06-09	58	63523.31	51295	30105.0	
1	1	166B05B0	1962-12-16	58	38175.46	15050	4995.0	
2	2	BFC13E88	1990-09-12	34	59863.77	24567	10323.0	
3	3	F2158F66	1985-11-03	35	84132.10	23712	3908.0	
4	4	7A73F334	1990-05-17	30	120512.00	93230	12158.4	

	fund_B	fund_C	fund_D	account_opened	last_transaction
0	4138.0	1420.0	15632.0	02-09-18	22-02-19
1	938.0	6696.0	2421.0	28-02-19	31-10-18
2	4590.0	8469.0	1185.0	25-04-18	02-04-18
3	492.0	6482.0	12830.0	07-11-17	08-11-18
4	51281.0	13434.0	18383.0	14-05-18	19-07-18

```
[38]: banking.dtypes
```

```
[38]:
```

Unnamed: 0	int64
cust_id	object
birth_date	object
Age	int64
acct_amount	float64
inv_amount	int64
fund_A	float64
fund_B	float64
fund_C	float64
fund_D	float64
account_opened	datetime64[ns]
last_transaction	object
acct_year	object
dtype:	object

```
[ ]: ##### Uniform dates
```

After having unified the currencies of different account amounts, I want to add
→ a temporal dimension to the analysis and see how customers have been
→ investing their money given the size of their account over each year. The
→ account_opened column represents when customers opened their accounts and is
→ a good proxy for segmenting customer activity and investment over time.

However, since this data was consolidated from multiple sources, I need to make sure that all dates are of the same format. I will do so by converting this column into a datetime object, while making sure that the format is inferred and potentially incorrect formats are set to missing.

- Printing the header of account_opened from the banking DataFrame and taking a look at the different results.

```
[30]: # Print the header of account_opened
print(banking['account_opened'].head())
```

```
0    02-09-18
1    28-02-19
2    25-04-18
3    07-11-17
4    14-05-18
Name: account_opened, dtype: object
```

Question

Take a look at the output. You tried converting the values to datetime using the default `to_datetime()` function without changing any argument, however received the following error:

ValueError: month must be in 1..12

Why do you think that is?

Possible Answers

The `to_datetime()` function needs to be explicitly told which date format each row is in.

The `to_datetime()` function can only be applied on YY-mm-dd date formats. (correct)

The 21-14-17 entry is erroneous and leads to an error.

- Converting the `account_opened` column to datetime, while making sure the date format is inferred and that erroneous formats that raise error return a missing value.

```
[31]: # Printing the header of account_opened
print(banking['account_opened'].head())

# Converting account_opened to datetime
banking['account_opened'] = pd.to_datetime(banking['account_opened'],
                                           # Inferring datetime format
                                           infer_datetime_format = True,
                                           # Returning missing value for error
                                           errors = 'coerce')
```

```
0    02-09-18
1    28-02-19
2    25-04-18
```

```
3    07-11-17
4    14-05-18
Name: account_opened, dtype: object
```

- Extracting the year from the amended account_opened column and assign it to the acct_year column.
- Printing the newly created acct_year column.

```
[32]: # Printing the header of account_opened
print(banking['account_opened'].head())

# Converting account_opened to datetime
banking['account_opened'] = pd.to_datetime(banking['account_opened'],
                                           # Infer datetime format
                                           infer_datetime_format = True,
                                           # Return missing value for error
                                           errors = 'coerce')

# Getting year of account opened
banking['acct_year'] = banking['account_opened'].dt.strftime('%Y')

# Printing acct_year
print(banking['acct_year'])
```

```
0    2018-02-09
1    2019-02-28
2    2018-04-25
3    2017-07-11
4    2018-05-14
Name: account_opened, dtype: datetime64[ns]
0      2018
1      2019
2      2018
3      2017
4      2018
...
95     2018
96     2017
97     2017
98     2017
99     2017
Name: acct_year, Length: 100, dtype: object
```

How's our data integrity?

New data has been merged into the banking DataFrame that contains details on how investments in the inv_amount column are allocated across four different funds A, B, C and D.

Furthermore, the age and birthdays of customers are now stored in the age and birth_date columns

respectively.

I want to understand how customers of different age groups invest. I want to first make sure the data I am analyzing is correct. I will do so by cross field checking values of `inv_amount` and `age` against the amount invested in different funds and customers' birthdays.

```
[33]: # Storing fund columns to sum against
fund_columns = ['fund_A', 'fund_B', 'fund_C', 'fund_D']

# Finding rows where fund_columns row sum == inv_amount
inv_equ = banking[fund_columns].sum(axis = 1) == banking['inv_amount']

# Storing consistent and inconsistent data
consistent_inv = banking[inv_equ]
inconsistent_inv = banking[~inv_equ]

# Storing consistent and inconsistent data
print("Number of inconsistent investments: ", inconsistent_inv.shape[0])
```

Number of inconsistent investments: 8

Storing today's date into `today`, and manually calculate customers' ages and storing them in `ages_manual`.

Finding all rows of `banking` where the `age` column is equal to `ages_manual` and then filtering `banking` into `consistent_ages` and `inconsistent_ages`.

```
[39]: # Converting birth_date to datetime
banking['birth_date'] = pd.to_datetime(banking['birth_date'],
                                     # Infer datetime format
                                     infer_datetime_format = True,
                                     # Return missing value for error
                                     errors = 'coerce')
```

```
[41]: # Storing today's date and find ages
today = dt.date.today()
ages_manual = today.year - banking['birth_date'].dt.year

# Finding rows where age column == ages_manual
age_equ = banking['Age'] == ages_manual

# Storing consistent and inconsistent data
consistent_ages = banking[age_equ]
inconsistent_ages = banking[~age_equ]

# Storing consistent and inconsistent data
print("Number of inconsistent ages: ", inconsistent_ages.shape[0])
```

Number of inconsistent ages: 8

Awesome work! There are only 8 and 8 rows affected by inconsistent `inv_amount` and `age` values respectively. In this case, it's best to investigate the underlying data sources before deciding on a course of action!

Is this missing at random?

There are a variety of missingness types when observing missing data. As a reminder, missingness types can be described as the following:

- Missing Completely at Random: No systematic relationship between a column's missing values and other or own values.
- Missing at Random: There is a systematic relationship between a column's missing values and other observed values.
- Missing not at Random: There is a systematic relationship between a column's missing values and unobserved values.

You have a DataFrame containing customer satisfaction scores for a service. What type of missingness is the following?

A customer `satisfaction_score` column with missing values for highly dissatisfied customers.

Missing completely at random

There's obviously an underlying mechanism behind the missingness of the data.

Missing at Random

This would be more applicable if missing `satisfaction_score` values were due to values in another column of the DataFrame.

Missing not at random

Awesome work! This is a clear example of missing not at random, where low values of `satisfaction_score` are missing because of inherently low satisfaction!

```
[45]: pip install missingno
```

Collecting missingno

Downloading <https://files.pythonhosted.org/packages/2b/de/6e4dd6d720c49939544352155dc06a08c9f7e4271aa631a559dfbeaaf9d4/missingno-0.4.2-py3-none-any.whl>

Requirement already satisfied: scipy in /opt/anaconda3/lib/python3.7/site-packages (from missingno) (1.3.1)

Requirement already satisfied: seaborn in /opt/anaconda3/lib/python3.7/site-packages (from missingno) (0.9.0)

Requirement already satisfied: matplotlib in /opt/anaconda3/lib/python3.7/site-packages (from missingno) (3.1.1)

Requirement already satisfied: numpy in /opt/anaconda3/lib/python3.7/site-packages (from missingno) (1.17.2)

Requirement already satisfied: pandas>=0.15.2 in /opt/anaconda3/lib/python3.7/site-packages (from seaborn->missingno) (0.25.1)


Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno) (0.10.0)

Requirement already satisfied: kiwisolver>=1.0.1 in
/opt/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno) (1.1.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in
/opt/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno) (2.4.2)
Requirement already satisfied: python-dateutil>=2.1 in
/opt/anaconda3/lib/python3.7/site-packages (from matplotlib->missingno) (2.8.0)
Requirement already satisfied: pytz>=2017.2 in
/opt/anaconda3/lib/python3.7/site-packages (from
pandas>=0.15.2->seaborn->missingno) (2019.3)
Requirement already satisfied: six in /opt/anaconda3/lib/python3.7/site-packages
(from cycler>=0.10->matplotlib->missingno) (1.12.0)
Requirement already satisfied: setuptools in /opt/anaconda3/lib/python3.7/site-
packages (from kiwisolver>=1.0.1->matplotlib->missingno) (41.4.0)
Installing collected packages: missingno
Successfully installed missingno-0.4.2
Note: you may need to restart the kernel to use updated packages.

```
[46]: importing matplotlib.pyplot as plt
import missingno as msno
# Print number of missing values in banking
print(banking.isna().sum())

# Visualizing missingness matrix
msno.matrix(banking)
plt.show()
```

```
Unnamed: 0      0
cust_id        0
birth_date      0
Age            0
acct_amount     0
inv_amount     0
fund_A         0
fund_B         0
fund_C         0
fund_D         0
account_opened  0
last_transaction 0
acct_year      0
dtype: int64
```




[47]: *# Isolating the values of banking missing values of inv_amount into missing_investors and with non-missing inv_amount values into investors.*

```
# Printing number of missing values in banking
print(banking.isna().sum())

# Visualizing missingness matrix
msno.matrix(banking)
plt.show()

# Isolating missing and non missing values of inv_amount
missing_investors = banking[banking['inv_amount'].isna()]
investors = banking[~banking['inv_amount'].isna()]
```

```
Unnamed: 0      0
cust_id         0
birth_date      0
Age             0
acct_amount     0
inv_amount      0
fund_A          0
fund_B          0
fund_C          0
fund_D          0
account_opened  0
last_transaction 0
acct_year       0
dtype: int64
```



That's all for now. I have used many other techniques for data cleaning in my other projects available on my Github page.

Thanks a lot for your attention.

[]: